

# 1. Import the MNIST Dataset.

```
In [2]: import tensorflow as tf  
import matplotlib.pyplot as plt
```


```
mnist = tf.keras.datasets.mnist  
(x_train, y_train),(x_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>


11493376/11490434 [=====] - 3s 0us/step

## 2. MNIST Dataset

```
In [3]: plt.imshow(x_train[0])  
plt.show()
```




```
In [4]: plt.imshow(x_train[0], cmap = plt.cm.binary)
plt.show()
```



### 3. Sequential Model - Normalize

```
In [5]: x_train = tf.keras.utils.normalize(x_train, axis=1)
x_test = tf.keras.utils.normalize(x_test, axis=1)
```

```
In [6]: plt.imshow(x_train[0], cmap = plt.cm.binary)
plt.show()
```



### 3. Input (32,32,10), SGD, batch\_size=64, epochs=20

```
In [7]: model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

from keras import optimizers

sgd = optimizers.SGD()
model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics
=['accuracy'])
model.fit(x_train, y_train, batch_size=64, epochs=20)
```

Using TensorFlow backend.

Train on 60000 samples

Epoch 1/20

60000/60000 [=====] - 1s 20us/sample - loss: 1.6825  
- accuracy: 0.5153

Epoch 2/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.7494  
- accuracy: 0.8041

Epoch 3/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.5016  
- accuracy: 0.8587

Epoch 4/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.4137  
- accuracy: 0.8834

Epoch 5/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.3693  
- accuracy: 0.8943

Epoch 6/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.3426  
- accuracy: 0.9026

Epoch 7/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.3236  
- accuracy: 0.9071

Epoch 8/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.3088  
- accuracy: 0.9108

Epoch 9/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.2962  
- accuracy: 0.9146

Epoch 10/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.2857  
- accuracy: 0.9172

Epoch 11/20

60000/60000 [=====] - 1s 21us/sample - loss: 0.2757  
- accuracy: 0.9199

Epoch 12/20

60000/60000 [=====] - 1s 24us/sample - loss: 0.2667  
- accuracy: 0.9227s - 1 - ETA: 0s - loss: 0.2655

Epoch 13/20

60000/60000 [=====] - ETA: 0s - loss: 0.2587 - accuracy: 0.9248  
ETA: 0s - loss: 0 - 1s 22us/sample - loss: 0.2582 - accuracy: 0.9250

Epoch 14/20

60000/60000 [=====] - 1s 21us/sample - loss: 0.2498  
- accuracy: 0.9273

Epoch 15/20

60000/60000 [=====] - 1s 22us/sample - loss: 0.2424  
- accuracy: 0.9291

Epoch 16/20

60000/60000 [=====] - 1s 22us/sample - loss: 0.2352  
- accuracy: 0.9317

Epoch 17/20

60000/60000 [=====] - ETA: 0s - loss: 0.2287 - accuracy: 0.93  
- 1s 23us/sample - loss: 0.2281 - accuracy: 0.9338

Epoch 18/20

60000/60000 [=====] - 1s 22us/sample - loss: 0.2216  
- accuracy: 0.9360s - loss: 0.2122 - accuracy

Epoch 19/20

```
60000/60000 [=====] - 1s 22us/sample - loss: 0.2152
- accuracy: 0.9378s - loss: 0.2
Epoch 20/20
60000/60000 [=====] - 1s 24us/sample - loss: 0.2093
- accuracy: 0.9397s - los
```

Out[7]: <tensorflow.python.keras.callbacks.History at 0x25aaff54748>

```
In [8]: val_loss, val_acc = model.evaluate(x_test, y_test)
print("Validation Accuracy: "+str(val_acc))
```

```
10000/10000 [=====] - 0s 24us/sample - loss: 0.2165
- accuracy: 0.9370
Validation Accuracy: 0.937
```

```
In [9]: predictions = model.predict([x_test])
```

```
In [10]: print(predictions)
```


```
[9.73212627e-06 1.25599300e-07 7.07117666e-04 ... 9.96481419e-01
 1.16260060e-06 7.62079144e-05]
[2.83717585e-04 4.52270004e-04 9.86150086e-01 ... 3.65136671e-10
 2.43056129e-04 2.51890973e-11]
[1.91952381e-07 9.97459948e-01 5.42321417e-04 ... 2.19743670e-04
 9.73697461e-04 3.27786875e-05]
...
[6.48240857e-08 1.88459962e-05 1.00796215e-05 ... 4.77399357e-04
 3.79088498e-03 1.18911592e-02]
[2.07536141e-04 4.87227917e-05 1.03505408e-06 ... 4.84856002e-07
 4.64795753e-02 3.91939602e-06]
[2.69137614e-04 1.29105308e-06 3.55077558e-03 ... 3.66304000e-08
 4.97732981e-05 9.02434323e-08]]
```

```
In [11]: import numpy as np

print(np.argmax([predictions[0]]))
```

7

```
In [12]: plt.imshow(x_test[0])  
plt.show()
```



## 4. Output(64,64,10)

When the number of neurons are increased, the time taken to train the model increases. The accuracy here didn't change significantly.

```
In [13]: model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(64, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(64, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

from keras import optimizers

sgd = optimizers.SGD()
model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics
=['accuracy'])
model.fit(x_train, y_train, batch_size=64, epochs=20)

val_loss, val_acc = model.evaluate(x_test, y_test)
print("Validation Accuracy: "+str(val_acc))
```



Train on 60000 samples

Epoch 1/20

60000/60000 [=====] - 1s 23us/sample - loss: 1.6257  
- accuracy: 0.5661

Epoch 2/20

60000/60000 [=====] - 1s 19us/sample - loss: 0.6221  
- accuracy: 0.8383

Epoch 3/20

60000/60000 [=====] - 1s 19us/sample - loss: 0.4409  
- accuracy: 0.8773

Epoch 4/20

60000/60000 [=====] - 1s 19us/sample - loss: 0.3764  
- accuracy: 0.8932

Epoch 5/20

60000/60000 [=====] - 1s 24us/sample - loss: 0.3415  
- accuracy: 0.9025

Epoch 6/20

60000/60000 [=====] - 2s 25us/sample - loss: 0.3184  
- accuracy: 0.9088

Epoch 7/20

60000/60000 [=====] - 2s 26us/sample - loss: 0.3005  
- accuracy: 0.9134s - loss: 0.3007 - accuracy

Epoch 8/20

60000/60000 [=====] - 2s 26us/sample - loss: 0.2858  
- accuracy: 0.9178s - loss:

Epoch 9/20

60000/60000 [=====] - 2s 26us/sample - loss: 0.2734  
- accuracy: 0.9217

Epoch 10/20

60000/60000 [=====] - 2s 27us/sample - loss: 0.2623  
- accuracy: 0.9253

Epoch 11/20

60000/60000 [=====] - 2s 27us/sample - loss: 0.2523  
- accuracy: 0.9283

Epoch 12/20

60000/60000 [=====] - 2s 25us/sample - loss: 0.2429  
- accuracy: 0.9310

Epoch 13/20

60000/60000 [=====] - 1s 25us/sample - loss: 0.2344  
- accuracy: 0.9328

Epoch 14/20

60000/60000 [=====] - 2s 25us/sample - loss: 0.2267  
- accuracy: 0.9350

Epoch 15/20

60000/60000 [=====] - 2s 26us/sample - loss: 0.2195  
- accuracy: 0.9371

Epoch 16/20

60000/60000 [=====] - 2s 25us/sample - loss: 0.2127  
- accuracy: 0.9390

Epoch 17/20

60000/60000 [=====] - 2s 26us/sample - loss: 0.2067  
- accuracy: 0.9407

Epoch 18/20

60000/60000 [=====] - 2s 25us/sample - loss: 0.2008  
- accuracy: 0.9427

Epoch 19/20

60000/60000 [=====] - 2s 25us/sample - loss: 0.1956

```
- accuracy: 0.9439
Epoch 20/20
60000/60000 [=====] - 2s 26us/sample - loss: 0.1903
- accuracy: 0.9456
10000/10000 [=====] - 0s 47us/sample - loss: 0.1958
- accuracy: 0.9425
Validation Accuracy: 0.9425
```

**5.Output(32,32,10), SGD(lr=0.01, decay\_rate=0.000001, momentum\_rate=0.9)**

```
In [14]: model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

from keras import optimizers

sgd = optimizers.SGD(lr=0.01, decay=0.000001, momentum=0.9)
model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics
=['accuracy'])
model.fit(x_train, y_train, batch_size=64, epochs=20)

val_loss, val_acc = model.evaluate(x_test, y_test)
print("Validation Accuracy: "+str(val_acc))
```

Train on 60000 samples

Epoch 1/20

60000/60000 [=====] - 1s 20us/sample - loss: 1.9127  
- accuracy: 0.4407

Epoch 2/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.8043  
- accuracy: 0.8040

Epoch 3/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.4888  
- accuracy: 0.8673

Epoch 4/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.4034  
- accuracy: 0.8861s -

Epoch 5/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.3644  
- accuracy: 0.8964

Epoch 6/20

60000/60000 [=====] - 1s 18us/sample - loss: 0.3402  
- accuracy: 0.9020

Epoch 7/20

60000/60000 [=====] - 1s 21us/sample - loss: 0.3226  
- accuracy: 0.9068

Epoch 8/20

60000/60000 [=====] - 1s 23us/sample - loss: 0.3085  
- accuracy: 0.9107s - loss: 0.3098 - accuracy

Epoch 9/20

60000/60000 [=====] - 1s 23us/sample - loss: 0.2966  
- accuracy: 0.9141

Epoch 10/20

60000/60000 [=====] - 1s 23us/sample - loss: 0.2857  
- accuracy: 0.9172

Epoch 11/20

60000/60000 [=====] - 1s 23us/sample - loss: 0.2759  
- accuracy: 0.9204

Epoch 12/20

60000/60000 [=====] - 1s 23us/sample - loss: 0.2668  
- accuracy: 0.9236s - loss: 0.2531 - accuracy - ETA:

Epoch 13/20

60000/60000 [=====] - 1s 23us/sample - loss: 0.2587  
- accuracy: 0.9252

Epoch 14/20

60000/60000 [=====] - 1s 24us/sample - loss: 0.2511  
- accuracy: 0.9280

Epoch 15/20

60000/60000 [=====] - 1s 25us/sample - loss: 0.2439  
- accuracy: 0.9300

Epoch 16/20

60000/60000 [=====] - 1s 23us/sample - loss: 0.2374  
- accuracy: 0.9318

Epoch 17/20

60000/60000 [=====] - 1s 23us/sample - loss: 0.2313  
- accuracy: 0.9335s - loss:

Epoch 18/20

60000/60000 [=====] - 1s 23us/sample - loss: 0.2256  
- accuracy: 0.9357

Epoch 19/20

60000/60000 [=====] - 1s 23us/sample - loss: 0.2200

```
- accuracy: 0.9374
Epoch 20/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.2151
- accuracy: 0.9390
10000/10000 [=====] - 0s 33us/sample - loss: 0.2154
- accuracy: 0.9354
Validation Accuracy: 0.9354
```

**Momentum is used for converging in the models. In our case, even if we use the momentum for faster convergence of gradients, there is no change in the accuracy. It is almost the same even without using it.**

**6. Input(32,32,10), batch\_size=128 and Input(32,32,10), batch\_size=32**

```
In [15]: model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

from keras import optimizers

sgd = optimizers.SGD(lr=0.01, decay=0.000001, momentum=0.9)
model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics
=['accuracy'])
model.fit(x_train, y_train, batch_size=128, epochs=20)

val_loss, val_acc = model.evaluate(x_test, y_test)
print("Validation Accuracy: "+str(val_acc))
```

Train on 60000 samples

Epoch 1/20

60000/60000 [=====] - 1s 13us/sample - loss: 2.1416  
- accuracy: 0.2872

Epoch 2/20

60000/60000 [=====] - 1s 11us/sample - loss: 1.4786  
- accuracy: 0.6686

Epoch 3/20

60000/60000 [=====] - 1s 10us/sample - loss: 0.8518  
- accuracy: 0.8022

Epoch 4/20

60000/60000 [=====] - 1s 11us/sample - loss: 0.6050  
- accuracy: 0.8416

Epoch 5/20

60000/60000 [=====] - 1s 11us/sample - loss: 0.5011  
- accuracy: 0.8615

Epoch 6/20

60000/60000 [=====] - 1s 12us/sample - loss: 0.4445  
- accuracy: 0.8748

Epoch 7/20

60000/60000 [=====] - 1s 14us/sample - loss: 0.4086  
- accuracy: 0.8834

Epoch 8/20

60000/60000 [=====] - 1s 14us/sample - loss: 0.3835  
- accuracy: 0.8901

Epoch 9/20

60000/60000 [=====] - 1s 14us/sample - loss: 0.3649  
- accuracy: 0.8953

Epoch 10/20

60000/60000 [=====] - 1s 13us/sample - loss: 0.3500  
- accuracy: 0.8995

Epoch 11/20

60000/60000 [=====] - 1s 14us/sample - loss: 0.3378  
- accuracy: 0.9025

Epoch 12/20

60000/60000 [=====] - 1s 14us/sample - loss: 0.3274  
- accuracy: 0.9053

Epoch 13/20

60000/60000 [=====] - 1s 14us/sample - loss: 0.3183  
- accuracy: 0.9089

Epoch 14/20

60000/60000 [=====] - 1s 13us/sample - loss: 0.3104  
- accuracy: 0.9108

Epoch 15/20

60000/60000 [=====] - 1s 14us/sample - loss: 0.3030  
- accuracy: 0.9128

Epoch 16/20

60000/60000 [=====] - 1s 13us/sample - loss: 0.2965  
- accuracy: 0.9145

Epoch 17/20

60000/60000 [=====] - 1s 13us/sample - loss: 0.2903  
- accuracy: 0.9169

Epoch 18/20

60000/60000 [=====] - 1s 14us/sample - loss: 0.2848  
- accuracy: 0.9188

Epoch 19/20

60000/60000 [=====] - 1s 14us/sample - loss: 0.2794

- accuracy: 0.9201

Epoch 20/20

60000/60000 [=====] - 1s 14us/sample - loss: 0.2744

- accuracy: 0.9216

10000/10000 [=====] - 0s 34us/sample - loss: 0.2696

- accuracy: 0.9253

Validation Accuracy: 0.9253



```
In [16]: model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

from keras import optimizers

sgd = optimizers.SGD(lr=0.01, decay=0.000001, momentum=0.9)
model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics
=['accuracy'])
model.fit(x_train, y_train, batch_size=32, epochs=20)

val_loss, val_acc = model.evaluate(x_test, y_test)
print("Validation Accuracy: "+str(val_acc))
```

Train on 60000 samples

Epoch 1/20

60000/60000 [=====] - 2s 38us/sample - loss: 1.3381  
- accuracy: 0.6321

Epoch 2/20

60000/60000 [=====] - 2s 38us/sample - loss: 0.4490  
- accuracy: 0.8742

Epoch 3/20

60000/60000 [=====] - 2s 38us/sample - loss: 0.3525  
- accuracy: 0.8989s - loss: 0

Epoch 4/20

60000/60000 [=====] - 2s 40us/sample - loss: 0.3148  
- accuracy: 0.9094s - loss: 0.3224 - - ETA: 0s - loss: 0.3

Epoch 5/20

60000/60000 [=====] - 2s 39us/sample - loss: 0.2903  
- accuracy: 0.9164s - los - ETA: 1s - loss: 0 - ETA

Epoch 6/20

60000/60000 [=====] - 2s 37us/sample - loss: 0.2716  
- accuracy: 0.9215

Epoch 7/20

60000/60000 [=====] - 2s 38us/sample - loss: 0.2567  
- accuracy: 0.9260s - loss: 0.2626 - accuracy: 0.92 - E

Epoch 8/20

60000/60000 [=====] - 2s 41us/sample - loss: 0.2435  
- accuracy: 0.9301

Epoch 9/20

60000/60000 [=====] - 2s 37us/sample - loss: 0.2318  
- accuracy: 0.9339

Epoch 10/20

60000/60000 [=====] - 2s 38us/sample - loss: 0.2219  
- accuracy: 0.9365

Epoch 11/20

60000/60000 [=====] - 2s 37us/sample - loss: 0.2131  
- accuracy: 0.9390

Epoch 12/20

60000/60000 [=====] - 2s 38us/sample - loss: 0.2051  
- accuracy: 0.9420

Epoch 13/20

60000/60000 [=====] - 2s 37us/sample - loss: 0.1977  
- accuracy: 0.9438

Epoch 14/20

60000/60000 [=====] - 2s 37us/sample - loss: 0.1909  
- accuracy: 0.9456

Epoch 15/20

60000/60000 [=====] - ETA: 0s - loss: 0.1844 - accur  
acy: 0.9472 ETA: 0s - loss: 0.1840 - accuracy - 2s 40us/sample - loss: 0.1845  
- accuracy: 0.9472

Epoch 16/20

60000/60000 [=====] - 2s 38us/sample - loss: 0.1784  
- accuracy: 0.9494s

Epoch 17/20

60000/60000 [=====] - 2s 37us/sample - loss: 0.1732  
- accuracy: 0.9509

Epoch 18/20

60000/60000 [=====] - 2s 37us/sample - loss: 0.1677  
- accuracy: 0.9525

Epoch 19/20

```
60000/60000 [=====] - 2s 38us/sample - loss: 0.1627
- accuracy: 0.9533
Epoch 20/20
60000/60000 [=====] - 2s 37us/sample - loss: 0.1578
- accuracy: 0.9551
10000/10000 [=====] - 0s 30us/sample - loss: 0.1703
- accuracy: 0.9503
Validation Accuracy: 0.9503
```

#conclusion Validation accuracy for batch\_size\_128 = 0.91 Validation accuracy for batch\_size\_32 = 0.95 Conclusion:  
There is a slight increase in accuracy as batch size is decreased.

## 7. Adding hidden layer when input (32,32,32,10)

```
In [17]: model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

from keras import optimizers

sgd = optimizers.SGD(lr=0.01, decay=0.000001, momentum=0.9)
model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics
=['accuracy'])
model.fit(x_train, y_train, batch_size=32, epochs=20)

val_loss, val_acc = model.evaluate(x_test, y_test)
print("Validation Accuracy: "+str(val_acc))
```

Train on 60000 samples

Epoch 1/20

60000/60000 [=====] - 2s 32us/sample - loss: 1.4222  
- accuracy: 0.6020

Epoch 2/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.4617  
- accuracy: 0.8675

Epoch 3/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.3505  
- accuracy: 0.8986

Epoch 4/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.3022  
- accuracy: 0.9120

Epoch 5/20

60000/60000 [=====] - 2s 31us/sample - loss: 0.2697  
- accuracy: 0.9205

Epoch 6/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.2439  
- accuracy: 0.9278

Epoch 7/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.2232  
- accuracy: 0.9338

Epoch 8/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.2058  
- accuracy: 0.9389

Epoch 9/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1918  
- accuracy: 0.9431

Epoch 10/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1795  
- accuracy: 0.9467

Epoch 11/20

60000/60000 [=====] - 2s 29us/sample - loss: 0.1687  
- accuracy: 0.9496

Epoch 12/20

60000/60000 [=====] - 2s 31us/sample - loss: 0.1587  
- accuracy: 0.9532

Epoch 13/20

60000/60000 [=====] - 2s 31us/sample - loss: 0.1506  
- accuracy: 0.9552

Epoch 14/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1424  
- accuracy: 0.9580

Epoch 15/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1359  
- accuracy: 0.9590

Epoch 16/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1296  
- accuracy: 0.9610

Epoch 17/20

60000/60000 [=====] - 2s 29us/sample - loss: 0.1243  
- accuracy: 0.9626

Epoch 18/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1187  
- accuracy: 0.9644

Epoch 19/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1140

```
- accuracy: 0.9658
Epoch 20/20
60000/60000 [=====] - 2s 30us/sample - loss: 0.1094
- accuracy: 0.9667
10000/10000 [=====] - 0s 23us/sample - loss: 0.1383
- accuracy: 0.9596
Validation Accuracy: 0.9596
```

**Conclusion: Addition of extra hidden layer has increased the processing time but did not bring any increase in accuracy.**

**8. Input (128,128,10) and adding Drop out layer where dropout = 0.5**

**Dropout is used to prevent overfitting.**

```
In [18]: model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dropout(rate=0.5, noise_shape=None, seed=None))
model.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
model.add(tf.keras.layers.Dropout(rate=0.5, noise_shape=None, seed=None))
model.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
model.add(tf.keras.layers.Dropout(rate=0.5, noise_shape=None, seed=None))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

from keras import optimizers

sgd = optimizers.SGD(lr=0.01, decay=0.000001, momentum=0.9)
model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics
=['accuracy'])
model.fit(x_train, y_train, batch_size=32, epochs=20)

val_loss, val_acc = model.evaluate(x_test, y_test)
print("Validation Accuracy: "+str(val_acc))
```

Train on 60000 samples

Epoch 1/20

60000/60000 [=====] - 3s 45us/sample - loss: 1.7499  
- accuracy: 0.3906

Epoch 2/20

60000/60000 [=====] - 3s 46us/sample - loss: 0.9973  
- accuracy: 0.6687

Epoch 3/20

60000/60000 [=====] - 3s 44us/sample - loss: 0.8054  
- accuracy: 0.7421

Epoch 4/20

60000/60000 [=====] - 3s 44us/sample - loss: 0.7122  
- accuracy: 0.7751

Epoch 5/20

60000/60000 [=====] - 3s 44us/sample - loss: 0.6479  
- accuracy: 0.7987

Epoch 6/20

60000/60000 [=====] - 3s 44us/sample - loss: 0.5990  
- accuracy: 0.8152

Epoch 7/20

60000/60000 [=====] - 3s 44us/sample - loss: 0.5715  
- accuracy: 0.8248

Epoch 8/20

60000/60000 [=====] - 3s 47us/sample - loss: 0.5434  
- accuracy: 0.8346

Epoch 9/20

60000/60000 [=====] - 3s 45us/sample - loss: 0.5199  
- accuracy: 0.8407

Epoch 10/20

60000/60000 [=====] - 3s 46us/sample - loss: 0.5018  
- accuracy: 0.8490

Epoch 11/20

60000/60000 [=====] - 3s 47us/sample - loss: 0.4809  
- accuracy: 0.8569

Epoch 12/20

60000/60000 [=====] - 3s 46us/sample - loss: 0.4666  
- accuracy: 0.8587

Epoch 13/20

60000/60000 [=====] - 3s 46us/sample - loss: 0.4558  
- accuracy: 0.8624

Epoch 14/20

60000/60000 [=====] - 3s 48us/sample - loss: 0.4420  
- accuracy: 0.8677

Epoch 15/20

60000/60000 [=====] - 3s 48us/sample - loss: 0.4307  
- accuracy: 0.8704

Epoch 16/20

60000/60000 [=====] - 3s 46us/sample - loss: 0.4200  
- accuracy: 0.8727

Epoch 17/20

60000/60000 [=====] - 3s 46us/sample - loss: 0.4088  
- accuracy: 0.8780

Epoch 18/20

60000/60000 [=====] - 3s 58us/sample - loss: 0.4014  
- accuracy: 0.8793

Epoch 19/20

60000/60000 [=====] - 4s 73us/sample - loss: 0.3974



```
- accuracy: 0.8795  
Epoch 20/20  
60000/60000 [=====] - 4s 63us/sample - loss: 0.3916  
- accuracy: 0.8821  
10000/10000 [=====] - 0s 39us/sample - loss: 0.1805  
- accuracy: 0.9452  
Validation Accuracy: 0.9452
```

**Adding the dropout layer prevents the overfitting which is good for the model not to memorize the patterns. Accuracy is not increased drastically.**

1. Model with better parameters

```
In [28]: model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten())
#model.add(tf.keras.layers.Dropout(rate=0.5, noise_shape=None, seed=None))
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
#model.add(tf.keras.layers.Dropout(rate=0.5, noise_shape=None, seed=None))
model.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
#model.add(tf.keras.layers.Dropout(rate=0.5, noise_shape=None, seed=None))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

from keras import optimizers

sgd = optimizers.SGD(lr=0.09, decay=0.000001, momentum=0.9)
model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics
=['accuracy'])
model.fit(x_train, y_train, batch_size=32, epochs=20)

val_loss, val_acc = model.evaluate(x_test, y_test)
print("Validation Accuracy: "+str(val_acc))
```

Train on 60000 samples

Epoch 1/20

60000/60000 [=====] - 2s 31us/sample - loss: 1.2566  
- accuracy: 0.6555

Epoch 2/20

60000/60000 [=====] - 2s 29us/sample - loss: 0.4504  
- accuracy: 0.8741

Epoch 3/20

60000/60000 [=====] - 2s 29us/sample - loss: 0.3582  
- accuracy: 0.8965

Epoch 4/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.3176  
- accuracy: 0.9080

Epoch 5/20

60000/60000 [=====] - 2s 31us/sample - loss: 0.2898  
- accuracy: 0.9156

Epoch 6/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.2677  
- accuracy: 0.9228

Epoch 7/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.2491  
- accuracy: 0.9277

Epoch 8/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.2336  
- accuracy: 0.9334

Epoch 9/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.2201  
- accuracy: 0.9365

Epoch 10/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.2085  
- accuracy: 0.9396

Epoch 11/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1982  
- accuracy: 0.9423

Epoch 12/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1892  
- accuracy: 0.9447

Epoch 13/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1811  
- accuracy: 0.9470

Epoch 14/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1733  
- accuracy: 0.9488

Epoch 15/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1669  
- accuracy: 0.9510

Epoch 16/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1607  
- accuracy: 0.9530

Epoch 17/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1550  
- accuracy: 0.9546

Epoch 18/20

60000/60000 [=====] - 2s 29us/sample - loss: 0.1499  
- accuracy: 0.9561

Epoch 19/20

60000/60000 [=====] - 2s 30us/sample - loss: 0.1449

```
- accuracy: 0.9574
Epoch 20/20
60000/60000 [=====] - 2s 29us/sample - loss: 0.1402
- accuracy: 0.9594
10000/10000 [=====] - 0s 23us/sample - loss: 0.1577
- accuracy: 0.9540
Validation Accuracy: 0.954
```

Conclusion: I think, for the given data, model with 2 hidden layers with 32,32 neurons with learning rate = 0.09 and batch size=32 will have good accuracy. Further increase in learning rate or hidden layers or batch size would decrease the accuracy progressively.

Question 10 -

Given the data, below parameters are important a. Number of hidden layers b. Number of neurons per layer c. Batch\_size d. Learning rate

In [ ]: