

EECS 738 Lab 4

Step 0: Import relevant packages

```
In [52]: from __future__ import print_function

import os

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import *
from tensorflow.keras.optimizers import SGD, Adam
from tensorflow.keras.activations import relu
from tensorflow.keras.regularizers import l2
from tensorflow.keras.constraints import max_norm
from tensorflow.keras import backend as K
from tensorflow.keras.datasets import mnist, cifar10
from keras.callbacks import ModelCheckpoint, LearningRateScheduler
from keras.callbacks import ReduceLROnPlateau

print('The Tensorflow version is {}'.format(tf.__version__))
print('The Keras version is {}'.format(keras.__version__))
print('The Pandas version is {}'.format(pd.__version__))

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

print("Packages Loaded")

def lr_schedule(epoch):
    """Learning Rate Schedule

    Learning rate is scheduled to be reduced after 80, 120, 160, 180 epochs.
    Called automatically every epoch as part of callbacks during training.

    # Arguments
        epoch (int): The number of epochs

    # Returns
        lr (float32): Learning rate
    """
```

```
lr = 1e-3
if epoch > 25:
    lr *= 0.5e-3
elif epoch > 20:
    lr *= 1e-3
elif epoch > 15:
    lr *= 1e-2
elif epoch > 10:
    lr *= 1e-1
print('Learning rate: ', lr)
return lr
```

The Tensorflow version is 2.1.0.

The Keras version is 2.2.4-tf.

The Pandas version is 0.23.4.

Packages Loaded

Step 1: Load data in to train and test splits

```
In [53]: # Lets import our MNIST data like we did in Lab 3.
mnist = tf.keras.datasets.mnist
cifar = tf.keras.datasets.cifar10
(mx_train, my_train), (mx_test, my_test) = mnist.load_data()
(cx_train, cy_train), (cx_test, cy_test) = cifar.load_data()

print('Shape of mnist train data {}'.format(mx_train.shape))
print('Shape of mnist test data {}'.format(mx_test.shape))

print('Shape of cifar train data {}'.format(cx_train.shape))
print('Shape of cifar test data {}'.format(cx_test.shape))

print(type(mx_test))
```

Shape of mnist train data (60000, 28, 28).

Shape of mnist test data (10000, 28, 28).

Shape of cifar train data (50000, 32, 32, 3).

Shape of cifar test data (10000, 32, 32, 3).

<class 'numpy.ndarray'>

Step2: Prepare the data

```
In [54]: #Normalize the data
mx_train = mx_train/255
mx_test = mx_test / 255

mx_train = tf.keras.utils.normalize(mx_train, axis=1)
mx_test = tf.keras.utils.normalize(mx_test, axis=1)

mx_train_hold = mx_train
my_train_hold = my_train

cx_train = tf.keras.utils.normalize(cx_train, axis=1)
cx_test = tf.keras.utils.normalize(cx_test, axis=1)

cx_train = cx_train.astype('float32')
cx_test = cx_test.astype('float32')
cx_train /= 255
cx_test /= 255

#im = mx_train[4,:,:]
#plt.imshow(im, cmap=plt.cm.binary)

#x_train = x_train.reshape(60000, 784)
#x_test = x_test.reshape(10000, 784)
#x_train = x_train.astype('float32')
#x_test = x_test.astype('float32')

#y_train = tf.keras.utils.to_categorical(y_train, 10)
#y_test = tf.keras.utils.to_categorical(y_test, 10)
# Now import Cifar-10 data and process it.
 #(train, target), (test, test_target) = cifar10.Load_data()
# Fill in the rest.
```

Step 3: Sequential NN shape input, 64, 64, 64, 10

```
In [4]: # Our baseline model for this Lab.
model = keras.models.Sequential()
model.add(keras.layers.Flatten(input_shape=[28,28]))

model.add(keras.layers.Dense(64, activation="relu"))
model.add(keras.layers.Dense(64, activation="relu"))
model.add(keras.layers.Dense(64, activation="relu"))
model.add(keras.layers.Dense(10, activation="softmax"))

model.compile(loss='sparse_categorical_crossentropy',
              optimizer=SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True),
              metrics=['accuracy'])

#model.fit(mx_train, my_train, epochs=20, batch_size=64, validation_data=(mx_test, my_test))
model_detail = model.fit(mx_train, my_train, epochs=20, batch_size=64, validation_split=0.1)

#plot accuracies for each epoch
history = pd.DataFrame(model_detail.history)

history.plot(figsize=(10,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.xlabel('Epochs')
plt.title('Sequential Neural Nets')
plt.show()

#After training the model, evaluate the test set
model.evaluate(mx_test,my_test)

#Print the summary of the model
model.summary()
```

Train on 54000 samples, validate on 6000 samples

Epoch 1/20

54000/54000 [=====] - 2s 32us/sample - loss: 0.5926 - accuracy: 0.8233 - val_loss: 0.2111 - val_accuracy: 0.9410

Epoch 2/20

54000/54000 [=====] - 1s 27us/sample - loss: 0.2171 - accuracy: 0.9357 - val_loss: 0.1543 - val_accuracy: 0.9538

Epoch 3/20

54000/54000 [=====] - 2s 28us/sample - loss: 0.1579 - accuracy: 0.9523 - val_loss: 0.1278 - val_accuracy: 0.9632

Epoch 4/20

54000/54000 [=====] - 1s 27us/sample - loss: 0.1261 - accuracy: 0.9624 - val_loss: 0.1137 - val_accuracy: 0.9665

Epoch 5/20

54000/54000 [=====] - 1s 28us/sample - loss: 0.1049 - accuracy: 0.9691 - val_loss: 0.1141 - val_accuracy: 0.9673

Epoch 6/20

54000/54000 [=====] - 1s 27us/sample - loss: 0.0888 - accuracy: 0.9729 - val_loss: 0.1021 - val_accuracy: 0.9700

Epoch 7/20

54000/54000 [=====] - 1s 27us/sample - loss: 0.0762 - accuracy: 0.9765 - val_loss: 0.0990 - val_accuracy: 0.9710

Epoch 8/20

54000/54000 [=====] - 1s 27us/sample - loss: 0.0661 - accuracy: 0.9797 - val_loss: 0.0988 - val_accuracy: 0.9710

Epoch 9/20

54000/54000 [=====] - 2s 28us/sample - loss: 0.0564 - accuracy: 0.9825 - val_loss: 0.0935 - val_accuracy: 0.9727

Epoch 10/20

54000/54000 [=====] - 2s 29us/sample - loss: 0.0514 - accuracy: 0.9841 - val_loss: 0.0952 - val_accuracy: 0.9722

Epoch 11/20

54000/54000 [=====] - 1s 28us/sample - loss: 0.0452 - accuracy: 0.9866 - val_loss: 0.0951 - val_accuracy: 0.9730

Epoch 12/20

54000/54000 [=====] - 2s 28us/sample - loss: 0.0400 - accuracy: 0.9867 - val_loss: 0.1031 - val_accuracy: 0.9703

Epoch 13/20

54000/54000 [=====] - 2s 28us/sample - loss: 0.0347 - accuracy: 0.9894 - val_loss: 0.1052 - val_accuracy: 0.9717

Epoch 14/20

54000/54000 [=====] - 1s 27us/sample - loss: 0.0302 - accuracy: 0.9907 - val_loss: 0.1083 - val_accuracy: 0.9698


```
Epoch 15/20
54000/54000 [=====] - 1s 27us/sample - loss: 0.0265 - accuracy: 0.9920 - val_loss:
0.1099 - val_accuracy: 0.9712
Epoch 16/20
54000/54000 [=====] - 2s 28us/sample - loss: 0.0237 - accuracy: 0.9926 - val_loss:
0.1136 - val_accuracy: 0.9700
Epoch 17/20
54000/54000 [=====] - 1s 27us/sample - loss: 0.0205 - accuracy: 0.9939 - val_loss:
0.1191 - val_accuracy: 0.9705
Epoch 18/20
54000/54000 [=====] - 1s 27us/sample - loss: 0.0174 - accuracy: 0.9950 - val_loss:
0.1213 - val_accuracy: 0.9718
Epoch 19/20
54000/54000 [=====] - 1s 27us/sample - loss: 0.0155 - accuracy: 0.9954 - val_loss:
0.1175 - val_accuracy: 0.9728
Epoch 20/20
54000/54000 [=====] - 2s 30us/sample - loss: 0.0131 - accuracy: 0.9966 - val_loss:
0.1219 - val_accuracy: 0.9728
```

Out[4]: <matplotlib.axes._subplots.AxesSubplot at 0x26dbce3feb8>

Out[4]: (0, 1)

Out[4]: Text(0.5, 0, 'Epochs')

Out[4]: Text(0.5, 1.0, 'Sequential Neural Nets')



10000/10000 [=====] - 0s 32us/sample - loss: 0.1254 - accuracy: 0.9689

Out[4]: [0.12535318766142883, 0.9689]

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 64)	50240
dense_1 (Dense)	(None, 64)	4160
dense_2 (Dense)	(None, 64)	4160
dense_3 (Dense)	(None, 10)	650
=====		
Total params: 59,210		
Trainable params: 59,210		
Non-trainable params: 0		

Step 4: Convert this baseline into a Functional Model using keras' Functional Model API.

<https://keras.io/getting-started/functional-api-guide/> (<https://keras.io/getting-started/functional-api-guide/>)

```
In [9]: # Create the functional Baseline here.
input_layer = keras.layers.Input(shape=mx_train.shape[1:])
il = Flatten()(input_layer)
h1 = keras.layers.Dense(64,activation="relu")(il)
h2 = keras.layers.Dense(64,activation="relu")(h1)
h3 = keras.layers.Dense(64,activation="relu")(h2)
output_layer = keras.layers.Dense(10, activation="softmax")(h3)
fmodel = keras.models.Model(inputs=[input_layer], outputs=[output_layer])

fmodel.compile(loss='sparse_categorical_crossentropy',
               optimizer=keras.optimizers.SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True),metrics=['accuracy'])

fmodel_detail = fmodel.fit(mx_train, my_train, epochs=20, batch_size=64, validation_split=0.1)

#plot accuracies for each epoch
fhistory = pd.DataFrame(fmodel_detail.history)
fhistory.plot(figsize=(10,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.xlabel('Epochs')
plt.title('Functional Neural Nets')
plt.show()

#After training the model, evaluate the test set
fmodel.evaluate(mx_test,my_test)

#Print the summary of the model
fmodel.summary()
```

Train on 54000 samples, validate on 6000 samples

Epoch 1/20

54000/54000 [=====] - 2s 31us/sample - loss: 0.6214 - accuracy: 0.8124 - val_loss: 0.2098 - val_accuracy: 0.9402

Epoch 2/20

54000/54000 [=====] - 2s 28us/sample - loss: 0.2175 - accuracy: 0.9365 - val_loss: 0.1518 - val_accuracy: 0.9555

Epoch 3/20

54000/54000 [=====] - 1s 28us/sample - loss: 0.1595 - accuracy: 0.9524 - val_loss: 0.1198 - val_accuracy: 0.9625

Epoch 4/20

54000/54000 [=====] - 1s 26us/sample - loss: 0.1259 - accuracy: 0.9619 - val_loss: 0.1091 - val_accuracy: 0.9670

Epoch 5/20

54000/54000 [=====] - 2s 30us/sample - loss: 0.1026 - accuracy: 0.9693 - val_loss: 0.1113 - val_accuracy: 0.9672

Epoch 6/20

54000/54000 [=====] - 2s 29us/sample - loss: 0.0883 - accuracy: 0.9734 - val_loss: 0.1065 - val_accuracy: 0.9700

Epoch 7/20

54000/54000 [=====] - 1s 26us/sample - loss: 0.0761 - accuracy: 0.9764 - val_loss: 0.0951 - val_accuracy: 0.9712

Epoch 8/20

54000/54000 [=====] - 1s 28us/sample - loss: 0.0651 - accuracy: 0.9801 - val_loss: 0.1056 - val_accuracy: 0.9680

Epoch 9/20

54000/54000 [=====] - 1s 28us/sample - loss: 0.0569 - accuracy: 0.9823 - val_loss: 0.0982 - val_accuracy: 0.9722

Epoch 10/20

54000/54000 [=====] - 1s 27us/sample - loss: 0.0497 - accuracy: 0.9850 - val_loss: 0.0925 - val_accuracy: 0.9742

Epoch 11/20

54000/54000 [=====] - 1s 27us/sample - loss: 0.0440 - accuracy: 0.9863 - val_loss: 0.0969 - val_accuracy: 0.9740

Epoch 12/20

54000/54000 [=====] - 2s 30us/sample - loss: 0.0389 - accuracy: 0.9879 - val_loss: 0.0936 - val_accuracy: 0.9757

Epoch 13/20

54000/54000 [=====] - 1s 28us/sample - loss: 0.0342 - accuracy: 0.9896 - val_loss: 0.0960 - val_accuracy: 0.9738

Epoch 14/20

54000/54000 [=====] - 1s 27us/sample - loss: 0.0298 - accuracy: 0.9909 - val_loss: 0.1034 - val_accuracy: 0.9745

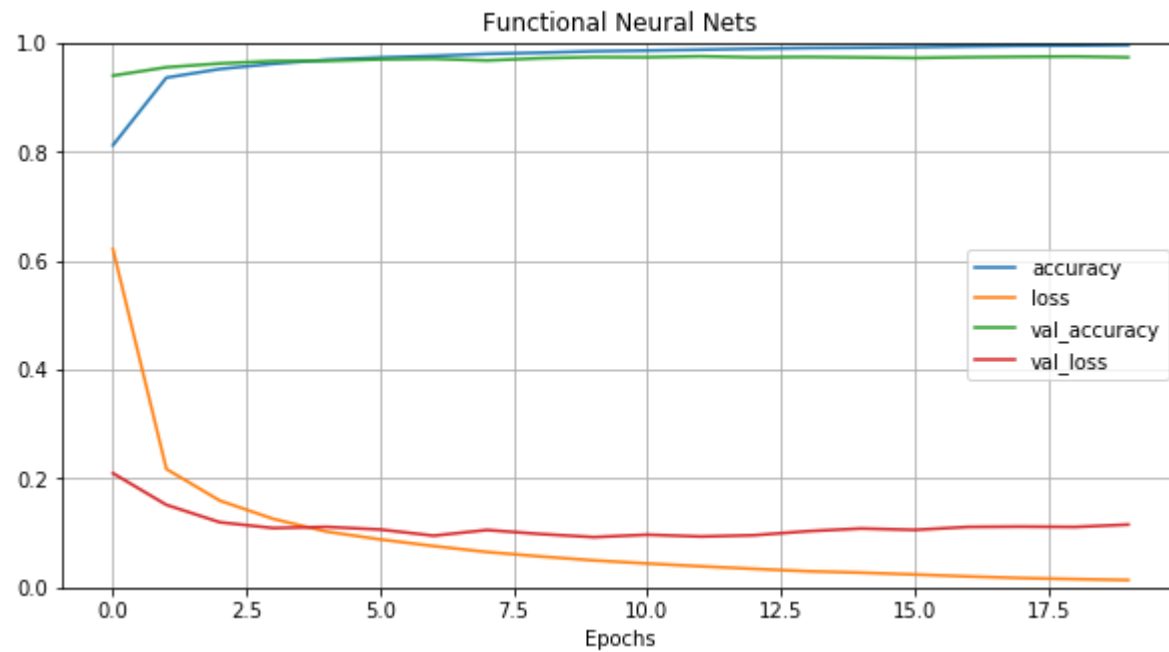
```
Epoch 15/20
54000/54000 [=====] - 1s 28us/sample - loss: 0.0272 - accuracy: 0.9915 - val_loss:
0.1083 - val_accuracy: 0.9737
Epoch 16/20
54000/54000 [=====] - 2s 28us/sample - loss: 0.0238 - accuracy: 0.9926 - val_loss:
0.1058 - val_accuracy: 0.9727
Epoch 17/20
54000/54000 [=====] - 1s 26us/sample - loss: 0.0203 - accuracy: 0.9939 - val_loss:
0.1112 - val_accuracy: 0.9740
Epoch 18/20
54000/54000 [=====] - 1s 27us/sample - loss: 0.0173 - accuracy: 0.9950 - val_loss:
0.1118 - val_accuracy: 0.9748
Epoch 19/20
54000/54000 [=====] - 1s 27us/sample - loss: 0.0154 - accuracy: 0.9954 - val_loss:
0.1111 - val_accuracy: 0.9753
Epoch 20/20
54000/54000 [=====] - 2s 28us/sample - loss: 0.0136 - accuracy: 0.9963 - val_loss:
0.1155 - val_accuracy: 0.9738
```

Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x26de10dba58>

Out[9]: (0, 1)

Out[9]: Text(0.5, 0, 'Epochs')

Out[9]: Text(0.5, 1.0, 'Functional Neural Nets')



10000/10000 [=====] - 0s 28us/sample - loss: 0.1253 - accuracy: 0.9690

Out[9]: [0.12528867619766387, 0.969]

Model: "model_1"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 784)]	0
=====		
flatten_2 (Flatten)	(None, 784)	0
=====		
dense_43 (Dense)	(None, 64)	50240
=====		
dense_44 (Dense)	(None, 64)	4160
=====		
dense_45 (Dense)	(None, 64)	4160
=====		
dense_46 (Dense)	(None, 10)	650
=====		
Total params: 59,210		
Trainable params: 59,210		
Non-trainable params: 0		
=====		

Step 5: Shallow ResNet for MNIST.

```
In [6]: inputs = tf.keras.Input(shape=(784,), name='img')
x = Dense(128, activation='relu')(inputs)
block_1_output = Dense(128, activation='relu')(x)

x = Dense(128)(block_1_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)

outputs = Dense(10, activation='softmax')(x)

rmodel = tf.keras.Model(inputs, outputs, name='resnet')

rmodel.compile(Adam(amsgrad=True), 'sparse_categorical_crossentropy', metrics=['accuracy'])

mx_train = mx_train.reshape(60000, 784)
mx_test = mx_test.reshape(10000, 784)
#mx_train = mx_train.astype('float32')
#mx_test = mx_test.astype('float32')

rmodel_detail = rmodel.fit(mx_train, my_train,
                           batch_size=128,
                           epochs=20,
                           validation_split=0.1)

#plot accuracies for each epoch
rhistory = pd.DataFrame(rmodel_detail.history)
rhistory.plot(figsize=(10,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.xlabel('Epochs')
plt.title('Residual Neural Nets')
plt.show()

#After training the model, evaluate the test set
rmodel.evaluate(mx_test, my_test)
```

```
#Print the summary of the model  
rmodel.summary()
```


Train on 54000 samples, validate on 6000 samples

Epoch 1/20

54000/54000 [=====] - 2s 45us/sample - loss: 0.5722 - accuracy: 0.8289 - val_loss: 0.1713 - val_accuracy: 0.9543

Epoch 2/20

54000/54000 [=====] - 2s 36us/sample - loss: 0.2086 - accuracy: 0.9440 - val_loss: 0.1236 - val_accuracy: 0.9645

Epoch 3/20

54000/54000 [=====] - 2s 37us/sample - loss: 0.1477 - accuracy: 0.9596 - val_loss: 0.1111 - val_accuracy: 0.9687

Epoch 4/20

54000/54000 [=====] - 2s 35us/sample - loss: 0.1118 - accuracy: 0.9696 - val_loss: 0.1039 - val_accuracy: 0.9728

Epoch 5/20

54000/54000 [=====] - 2s 34us/sample - loss: 0.0938 - accuracy: 0.9745 - val_loss: 0.0951 - val_accuracy: 0.9752

Epoch 6/20

54000/54000 [=====] - 2s 35us/sample - loss: 0.0754 - accuracy: 0.9796 - val_loss: 0.0917 - val_accuracy: 0.9767

Epoch 7/20

54000/54000 [=====] - 2s 35us/sample - loss: 0.0636 - accuracy: 0.9829 - val_loss: 0.0865 - val_accuracy: 0.9770

Epoch 8/20

54000/54000 [=====] - 2s 38us/sample - loss: 0.0510 - accuracy: 0.9852 - val_loss: 0.0865 - val_accuracy: 0.9777

Epoch 9/20

54000/54000 [=====] - 2s 36us/sample - loss: 0.0434 - accuracy: 0.9878 - val_loss: 0.0967 - val_accuracy: 0.9763

Epoch 10/20

54000/54000 [=====] - 2s 36us/sample - loss: 0.0412 - accuracy: 0.9887 - val_loss: 0.0955 - val_accuracy: 0.9778

Epoch 11/20

54000/54000 [=====] - 2s 35us/sample - loss: 0.0363 - accuracy: 0.9897 - val_loss: 0.1028 - val_accuracy: 0.9780

Epoch 12/20

54000/54000 [=====] - 2s 36us/sample - loss: 0.0321 - accuracy: 0.9906 - val_loss: 0.1033 - val_accuracy: 0.9748

Epoch 13/20

54000/54000 [=====] - 2s 36us/sample - loss: 0.0264 - accuracy: 0.9926 - val_loss: 0.0940 - val_accuracy: 0.9787

Epoch 14/20

54000/54000 [=====] - 2s 36us/sample - loss: 0.0239 - accuracy: 0.9930 - val_loss: 0.1132 - val_accuracy: 0.9753


```
Epoch 15/20
54000/54000 [=====] - 2s 37us/sample - loss: 0.0207 - accuracy: 0.9942 - val_loss:
0.0963 - val_accuracy: 0.9758
Epoch 16/20
54000/54000 [=====] - 2s 38us/sample - loss: 0.0208 - accuracy: 0.9941 - val_loss:
0.1051 - val_accuracy: 0.9780
Epoch 17/20
54000/54000 [=====] - 2s 37us/sample - loss: 0.0193 - accuracy: 0.9947 - val_loss:
0.1013 - val_accuracy: 0.9770
Epoch 18/20
54000/54000 [=====] - 2s 35us/sample - loss: 0.0175 - accuracy: 0.9950 - val_loss:
0.0998 - val_accuracy: 0.9783
Epoch 19/20
54000/54000 [=====] - 2s 35us/sample - loss: 0.0107 - accuracy: 0.9972 - val_loss:
0.0914 - val_accuracy: 0.9813
Epoch 20/20
54000/54000 [=====] - 2s 36us/sample - loss: 0.0133 - accuracy: 0.9960 - val_loss:
0.1152 - val_accuracy: 0.9768
```

Out[6]: <matplotlib.axes._subplots.AxesSubplot at 0x26dbf270b00>

Out[6]: (0, 1)

Out[6]: Text(0.5, 0, 'Epochs')

Out[6]: Text(0.5, 1.0, 'Residual Neural Nets')



10000/10000 [=====] - 0s 33us/sample - loss: 0.1320 - accuracy: 0.9724

Out[6]: [0.132009984324011, 0.9724]

Model: "resnet"

Layer (type)	Output Shape	Param #
=====	=====	=====
img (InputLayer)	[(None, 784)]	0
dense_8 (Dense)	(None, 128)	100480
dense_9 (Dense)	(None, 128)	16512
dense_10 (Dense)	(None, 128)	16512
batch_normalization (Batch Normalization)	(None, 128)	512
activation (Activation)	(None, 128)	0
dropout (Dropout)	(None, 128)	0
dense_11 (Dense)	(None, 128)	16512
batch_normalization_1 (Batch Normalization)	(None, 128)	512
activation_1 (Activation)	(None, 128)	0
dropout_1 (Dropout)	(None, 128)	0
dense_12 (Dense)	(None, 10)	1290
=====	=====	=====
Total params: 152,330		
Trainable params: 151,818		
Non-trainable params: 512		

Analysis:

- Accuracy: There is an increase in the training and validation accuracies.
- Speed: There has been slight increase in training speeds. Could be because of the complex nature of ResNets

Step 6: Now lets make a deeper ResNet. Make A network with 10 Residual Blocks.

```
In [7]: # Create the deep ResNet here.
inputs = tf.keras.Input(shape=(784,), name='img')
x = Dense(128, activation='relu')(inputs)
block_1_output = Dense(128, activation='relu')(x)

x = Dense(128)(block_1_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
block_2_output = tf.keras.layers.add([x, block_1_output])

# We will repeat above for as many times as we think our computer can handle. For now Lets make just three Residual Blocks.
x = Dense(128)(block_2_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
block_3_output = tf.keras.layers.add([x, block_2_output])

x = Dense(128)(block_3_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
block_4_output = tf.keras.layers.add([x, block_3_output])

x = Dense(128)(block_4_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
```

```
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
block_5_output = tf.keras.layers.add([x, block_4_output])

x = Dense(128)(block_5_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
block_6_output = tf.keras.layers.add([x, block_5_output])

x = Dense(128)(block_6_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
block_7_output = tf.keras.layers.add([x, block_6_output])

x = Dense(128)(block_7_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
block_8_output = tf.keras.layers.add([x, block_7_output])

x = Dense(128)(block_8_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
```

```
block_9_output = tf.keras.layers.add([x, block_8_output])

x = Dense(128)(block_9_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
block_10_output = tf.keras.layers.add([x, block_9_output])

# Now we cast the final Residual output into a dense layer to be able to classify the output easier.
x = Dense(128, activation='relu')(block_10_output)
x = Dropout(0.5)(x)

outputs = Dense(10, activation='softmax')(x)

r10model = tf.keras.Model(inputs, outputs, name='resnet')

r10model.compile(Adam(amsgrad=True), 'sparse_categorical_crossentropy', metrics=['accuracy'])

mx_train = mx_train.reshape(60000, 784)
mx_test = mx_test.reshape(10000, 784)
#mx_train = mx_train.astype('float32')
#mx_test = mx_test.astype('float32')

r10model_detail = r10model.fit(mx_train, my_train,
                               batch_size=128,
                               epochs=20,
                               validation_split=0.1)

#plot accuracies for each epoch
r10history = pd.DataFrame(r10model_detail.history)
r10history.plot(figsize=(10,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.xlabel('Epochs')
plt.title('Residual Neural Nets')
plt.show()

#After training the model, evaluate the test set
r10model.evaluate(mx_test,my_test)
```



```
#Print the summary of the model  
r10model.summary()
```

Train on 54000 samples, validate on 6000 samples

Epoch 1/20

54000/54000 [=====] - 12s 231us/sample - loss: 2.1663 - accuracy: 0.2324 - val_loss: 1.7797 - val_accuracy: 0.2842

Epoch 2/20

54000/54000 [=====] - 10s 176us/sample - loss: 1.1125 - accuracy: 0.6068 - val_loss: 0.4452 - val_accuracy: 0.8578

Epoch 3/20

54000/54000 [=====] - 10s 178us/sample - loss: 0.5670 - accuracy: 0.8381 - val_loss: 0.2244 - val_accuracy: 0.9340

Epoch 4/20

54000/54000 [=====] - 10s 176us/sample - loss: 0.3522 - accuracy: 0.9106 - val_loss: 0.1545 - val_accuracy: 0.9572

Epoch 5/20

54000/54000 [=====] - 10s 177us/sample - loss: 0.2654 - accuracy: 0.9353 - val_loss: 0.1361 - val_accuracy: 0.9613

Epoch 6/20

54000/54000 [=====] - 9s 174us/sample - loss: 0.2094 - accuracy: 0.9476 - val_loss: 0.1236 - val_accuracy: 0.9655

Epoch 7/20

54000/54000 [=====] - 9s 168us/sample - loss: 0.1783 - accuracy: 0.9543 - val_loss: 0.1064 - val_accuracy: 0.9708

Epoch 8/20

54000/54000 [=====] - 9s 162us/sample - loss: 0.1516 - accuracy: 0.9612 - val_loss: 0.1044 - val_accuracy: 0.9725

Epoch 9/20

54000/54000 [=====] - 9s 162us/sample - loss: 0.1353 - accuracy: 0.9652 - val_loss: 0.1006 - val_accuracy: 0.9722

Epoch 10/20

54000/54000 [=====] - 9s 165us/sample - loss: 0.1180 - accuracy: 0.9698 - val_loss: 0.0970 - val_accuracy: 0.9733

Epoch 11/20

54000/54000 [=====] - 9s 158us/sample - loss: 0.1073 - accuracy: 0.9724 - val_loss: 0.0950 - val_accuracy: 0.9742

Epoch 12/20

54000/54000 [=====] - 9s 169us/sample - loss: 0.0985 - accuracy: 0.9749 - val_loss: 0.0955 - val_accuracy: 0.9762

Epoch 13/20

54000/54000 [=====] - 8s 154us/sample - loss: 0.0851 - accuracy: 0.9774 - val_loss: 0.1015 - val_accuracy: 0.9737

Epoch 14/20

54000/54000 [=====] - 9s 170us/sample - loss: 0.0826 - accuracy: 0.9780 - val_loss: 0.0949 - val_accuracy: 0.9780


```
Epoch 15/20
54000/54000 [=====] - 8s 154us/sample - loss: 0.0818 - accuracy: 0.9809 - val_loss:
0.1946 - val_accuracy: 0.9612
Epoch 16/20
54000/54000 [=====] - 9s 168us/sample - loss: 0.0805 - accuracy: 0.9787 - val_loss:
0.0997 - val_accuracy: 0.9767
Epoch 17/20
54000/54000 [=====] - 9s 159us/sample - loss: 0.0626 - accuracy: 0.9838 - val_loss:
0.1005 - val_accuracy: 0.9762
Epoch 18/20
54000/54000 [=====] - 9s 163us/sample - loss: 0.0566 - accuracy: 0.9850 - val_loss:
0.0998 - val_accuracy: 0.9777
Epoch 19/20
54000/54000 [=====] - 9s 161us/sample - loss: 0.0501 - accuracy: 0.9861 - val_loss:
0.1035 - val_accuracy: 0.9755
Epoch 20/20
54000/54000 [=====] - 9s 161us/sample - loss: 0.0498 - accuracy: 0.9864 - val_loss:
0.1094 - val_accuracy: 0.9757
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x26dd9ca9b70>
```

```
Out[7]: (0, 1)
```

```
Out[7]: Text(0.5, 0, 'Epochs')
```

```
Out[7]: Text(0.5, 1.0, 'Residual Neural Nets')
```



10000/10000 [=====] - 1s 64us/sample - loss: 0.1344 - accuracy: 0.9719

Out[7]: [0.13437970478771605, 0.9719]

Model: "resnet"

Layer (type)	Output Shape	Param #	Connected to
img (InputLayer)	[(None, 784)]	0	
dense_13 (Dense)	(None, 128)	100480	img[0][0]
dense_14 (Dense)	(None, 128)	16512	dense_13[0][0]
dense_15 (Dense)	(None, 128)	16512	dense_14[0][0]
batch_normalization_2 (BatchNor	(None, 128)	512	dense_15[0][0]
activation_2 (Activation)	(None, 128)	0	batch_normalization_2[0][0]
dropout_2 (Dropout)	(None, 128)	0	activation_2[0][0]
dense_16 (Dense)	(None, 128)	16512	dropout_2[0][0]
batch_normalization_3 (BatchNor	(None, 128)	512	dense_16[0][0]
activation_3 (Activation)	(None, 128)	0	batch_normalization_3[0][0]
dropout_3 (Dropout)	(None, 128)	0	activation_3[0][0]
add (Add)	(None, 128)	0	dropout_3[0][0] dense_14[0][0]
dense_17 (Dense)	(None, 128)	16512	add[0][0]
batch_normalization_4 (BatchNor	(None, 128)	512	dense_17[0][0]
activation_4 (Activation)	(None, 128)	0	batch_normalization_4[0][0]
dropout_4 (Dropout)	(None, 128)	0	activation_4[0][0]
dense_18 (Dense)	(None, 128)	16512	dropout_4[0][0]
batch_normalization_5 (BatchNor	(None, 128)	512	dense_18[0][0]
activation_5 (Activation)	(None, 128)	0	batch_normalization_5[0][0]

dropout_5 (Dropout)	(None, 128)	0	activation_5[0][0]
add_1 (Add)	(None, 128)	0	dropout_5[0][0] add[0][0]
dense_19 (Dense)	(None, 128)	16512	add_1[0][0]
batch_normalization_6 (BatchNor	(None, 128)	512	dense_19[0][0]
activation_6 (Activation)	(None, 128)	0	batch_normalization_6[0][0]
dropout_6 (Dropout)	(None, 128)	0	activation_6[0][0]
dense_20 (Dense)	(None, 128)	16512	dropout_6[0][0]
batch_normalization_7 (BatchNor	(None, 128)	512	dense_20[0][0]
activation_7 (Activation)	(None, 128)	0	batch_normalization_7[0][0]
dropout_7 (Dropout)	(None, 128)	0	activation_7[0][0]
add_2 (Add)	(None, 128)	0	dropout_7[0][0] add_1[0][0]
dense_21 (Dense)	(None, 128)	16512	add_2[0][0]
batch_normalization_8 (BatchNor	(None, 128)	512	dense_21[0][0]
activation_8 (Activation)	(None, 128)	0	batch_normalization_8[0][0]
dropout_8 (Dropout)	(None, 128)	0	activation_8[0][0]
dense_22 (Dense)	(None, 128)	16512	dropout_8[0][0]
batch_normalization_9 (BatchNor	(None, 128)	512	dense_22[0][0]
activation_9 (Activation)	(None, 128)	0	batch_normalization_9[0][0]
dropout_9 (Dropout)	(None, 128)	0	activation_9[0][0]
add_3 (Add)	(None, 128)	0	dropout_9[0][0] add_2[0][0]

dense_23 (Dense)	(None, 128)	16512	add_3[0][0]
batch_normalization_10 (Batch Normalization)	(None, 128)	512	dense_23[0][0]
activation_10 (Activation)	(None, 128)	0	batch_normalization_10[0][0]
dropout_10 (Dropout)	(None, 128)	0	activation_10[0][0]
dense_24 (Dense)	(None, 128)	16512	dropout_10[0][0]
batch_normalization_11 (Batch Normalization)	(None, 128)	512	dense_24[0][0]
activation_11 (Activation)	(None, 128)	0	batch_normalization_11[0][0]
dropout_11 (Dropout)	(None, 128)	0	activation_11[0][0]
add_4 (Add)	(None, 128)	0	dropout_11[0][0] add_3[0][0]
dense_25 (Dense)	(None, 128)	16512	add_4[0][0]
batch_normalization_12 (Batch Normalization)	(None, 128)	512	dense_25[0][0]
activation_12 (Activation)	(None, 128)	0	batch_normalization_12[0][0]
dropout_12 (Dropout)	(None, 128)	0	activation_12[0][0]
dense_26 (Dense)	(None, 128)	16512	dropout_12[0][0]
batch_normalization_13 (Batch Normalization)	(None, 128)	512	dense_26[0][0]
activation_13 (Activation)	(None, 128)	0	batch_normalization_13[0][0]
dropout_13 (Dropout)	(None, 128)	0	activation_13[0][0]
add_5 (Add)	(None, 128)	0	dropout_13[0][0] add_4[0][0]
dense_27 (Dense)	(None, 128)	16512	add_5[0][0]
batch_normalization_14 (Batch Normalization)	(None, 128)	512	dense_27[0][0]
activation_14 (Activation)	(None, 128)	0	batch_normalization_14[0][0]

dropout_14 (Dropout)	(None, 128)	0	activation_14[0][0]
dense_28 (Dense)	(None, 128)	16512	dropout_14[0][0]
batch_normalization_15 (BatchNo	(None, 128)	512	dense_28[0][0]
activation_15 (Activation)	(None, 128)	0	batch_normalization_15[0][0]
dropout_15 (Dropout)	(None, 128)	0	activation_15[0][0]
add_6 (Add)	(None, 128)	0	dropout_15[0][0] add_5[0][0]
dense_29 (Dense)	(None, 128)	16512	add_6[0][0]
batch_normalization_16 (BatchNo	(None, 128)	512	dense_29[0][0]
activation_16 (Activation)	(None, 128)	0	batch_normalization_16[0][0]
dropout_16 (Dropout)	(None, 128)	0	activation_16[0][0]
dense_30 (Dense)	(None, 128)	16512	dropout_16[0][0]
batch_normalization_17 (BatchNo	(None, 128)	512	dense_30[0][0]
activation_17 (Activation)	(None, 128)	0	batch_normalization_17[0][0]
dropout_17 (Dropout)	(None, 128)	0	activation_17[0][0]
add_7 (Add)	(None, 128)	0	dropout_17[0][0] add_6[0][0]
dense_31 (Dense)	(None, 128)	16512	add_7[0][0]
batch_normalization_18 (BatchNo	(None, 128)	512	dense_31[0][0]
activation_18 (Activation)	(None, 128)	0	batch_normalization_18[0][0]
dropout_18 (Dropout)	(None, 128)	0	activation_18[0][0]
dense_32 (Dense)	(None, 128)	16512	dropout_18[0][0]

batch_normalization_19 (BatchNo (None, 128)		512	dense_32[0][0]
activation_19 (Activation)	(None, 128)	0	batch_normalization_19[0][0]
dropout_19 (Dropout)	(None, 128)	0	activation_19[0][0]
add_8 (Add)	(None, 128)	0	dropout_19[0][0] add_7[0][0]
dense_33 (Dense)	(None, 128)	16512	add_8[0][0]
dropout_20 (Dropout)	(None, 128)	0	dense_33[0][0]
dense_34 (Dense)	(None, 10)	1290	dropout_20[0][0]
=====			
Total params: 441,226			
Trainable params: 436,618			
Non-trainable params: 4,608			

Analysis:

- Compared to Shallow net, there is no increase in the Accuracy, but training time is increased significantly.
- We can see, in summary, significant increase in number of parameters tuned.
- I think, we do not need any extra hidden layers to handle the complexity of MNIST data. Shallow ResNet would do better with less training time.

Step 7: ResNets that have one and three skip connections.

```

In [8]: inputs = tf.keras.Input(shape=(784,), name='img')
x = Dense(128, activation='relu')(inputs)
block_1_output = Dense(128, activation='relu')(x)

# Here is a one skip connection block.
x = Dense(128)(block_1_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
block_2_output = tf.keras.layers.add([x, block_1_output])

# Here is a three skip connection block.
x = Dense(128)(block_2_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
block_3_output = tf.keras.layers.add([x, block_2_output])

x = Dense(128, activation='relu')(block_3_output)
x = Dropout(0.5)(x)
outputs = Dense(10, activation='softmax')(x)

r13model = tf.keras.Model(inputs, outputs, name='resnet')

r13model.compile(Adam(amsgrad=True), 'sparse_categorical_crossentropy', metrics=['accuracy'])

mx_train = mx_train.reshape(60000, 784)
mx_test = mx_test.reshape(10000, 784)
#mx_train = mx_train.astype('float32')
#mx_test = mx_test.astype('float32')

r13model_detail = r13model.fit(mx_train, my_train,
                                batch_size=128,
                                epochs=20,

```

```
validation_split=0.1)

#plot accuracies for each epoch
r13history = pd.DataFrame(r13model_detail.history)
r13history.plot(figsize=(10,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.xlabel('Epochs')
plt.title('Residual Neural Nets')
plt.show()

#After training the model, evaluate the test set
r13model.evaluate(mx_test,my_test)

#Print the summary of the model
r13model.summary()
```

Train on 54000 samples, validate on 6000 samples

Epoch 1/20

54000/54000 [=====] - 4s 72us/sample - loss: 0.8241 - accuracy: 0.7452 - val_loss: 0.1998 - val_accuracy: 0.9430

Epoch 2/20

54000/54000 [=====] - 3s 58us/sample - loss: 0.2604 - accuracy: 0.9300 - val_loss: 0.1220 - val_accuracy: 0.9650

Epoch 3/20

54000/54000 [=====] - 3s 58us/sample - loss: 0.1819 - accuracy: 0.9525 - val_loss: 0.1201 - val_accuracy: 0.9653

Epoch 4/20

54000/54000 [=====] - 3s 58us/sample - loss: 0.1338 - accuracy: 0.9639 - val_loss: 0.1097 - val_accuracy: 0.9670

Epoch 5/20

54000/54000 [=====] - 3s 57us/sample - loss: 0.1023 - accuracy: 0.9718 - val_loss: 0.0931 - val_accuracy: 0.9735

Epoch 6/20

54000/54000 [=====] - 3s 57us/sample - loss: 0.0892 - accuracy: 0.9759 - val_loss: 0.0998 - val_accuracy: 0.9738

Epoch 7/20

54000/54000 [=====] - 3s 57us/sample - loss: 0.0729 - accuracy: 0.9808 - val_loss: 0.0976 - val_accuracy: 0.9750

Epoch 8/20

54000/54000 [=====] - 3s 58us/sample - loss: 0.0609 - accuracy: 0.9821 - val_loss: 0.1003 - val_accuracy: 0.9758

Epoch 9/20

54000/54000 [=====] - 3s 59us/sample - loss: 0.0543 - accuracy: 0.9854 - val_loss: 0.1009 - val_accuracy: 0.9757

Epoch 10/20

54000/54000 [=====] - 3s 58us/sample - loss: 0.0444 - accuracy: 0.9875 - val_loss: 0.0961 - val_accuracy: 0.9783

Epoch 11/20

54000/54000 [=====] - 3s 58us/sample - loss: 0.0395 - accuracy: 0.9886 - val_loss: 0.1004 - val_accuracy: 0.9765

Epoch 12/20

54000/54000 [=====] - 3s 57us/sample - loss: 0.0340 - accuracy: 0.9902 - val_loss: 0.0931 - val_accuracy: 0.9788

Epoch 13/20

54000/54000 [=====] - 3s 58us/sample - loss: 0.0301 - accuracy: 0.9914 - val_loss: 0.0988 - val_accuracy: 0.9767

Epoch 14/20

54000/54000 [=====] - 3s 58us/sample - loss: 0.0289 - accuracy: 0.9915 - val_loss: 0.1064 - val_accuracy: 0.9760

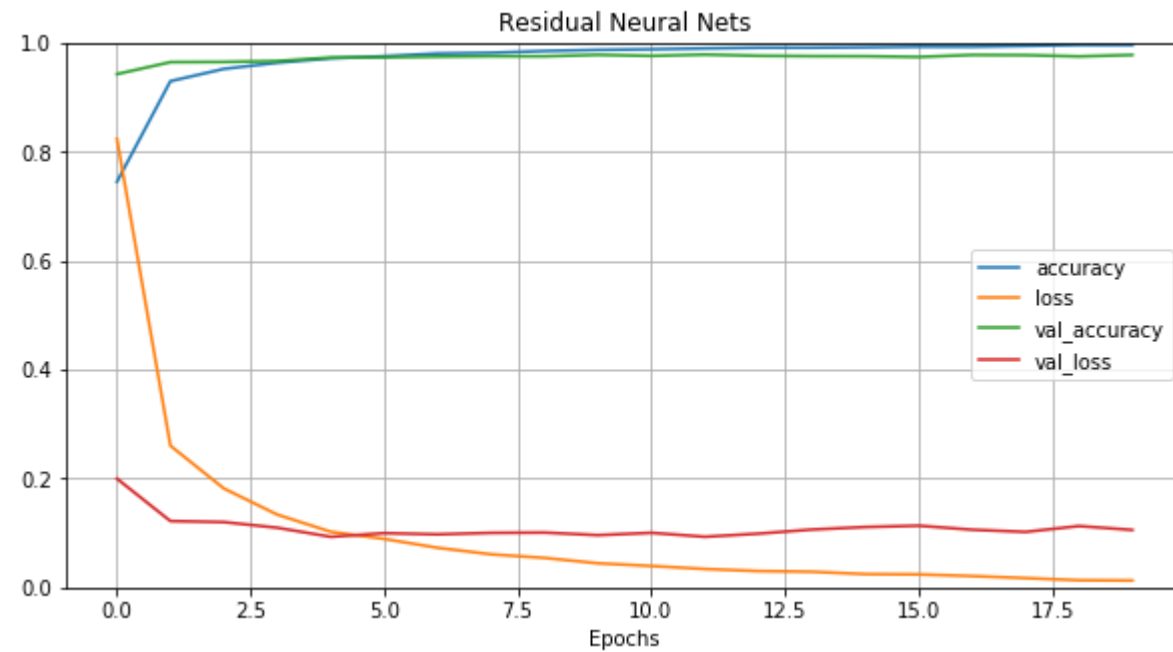
```
Epoch 15/20
54000/54000 [=====] - 3s 57us/sample - loss: 0.0246 - accuracy: 0.9923 - val_loss:
0.1110 - val_accuracy: 0.9758
Epoch 16/20
54000/54000 [=====] - 3s 57us/sample - loss: 0.0240 - accuracy: 0.9933 - val_loss:
0.1135 - val_accuracy: 0.9745
Epoch 17/20
54000/54000 [=====] - 3s 57us/sample - loss: 0.0210 - accuracy: 0.9933 - val_loss:
0.1061 - val_accuracy: 0.9782
Epoch 18/20
54000/54000 [=====] - 3s 59us/sample - loss: 0.0172 - accuracy: 0.9947 - val_loss:
0.1021 - val_accuracy: 0.9778
Epoch 19/20
54000/54000 [=====] - 3s 59us/sample - loss: 0.0133 - accuracy: 0.9965 - val_loss:
0.1128 - val_accuracy: 0.9755
Epoch 20/20
54000/54000 [=====] - 3s 58us/sample - loss: 0.0128 - accuracy: 0.9964 - val_loss:
0.1057 - val_accuracy: 0.9782
```

Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x26ddefe75c0>

Out[8]: (0, 1)

Out[8]: Text(0.5, 0, 'Epochs')

Out[8]: Text(0.5, 1.0, 'Residual Neural Nets')



10000/10000 [=====] - 0s 40us/sample - loss: 0.1323 - accuracy: 0.9757

Out[8]: [0.13233828188165084, 0.9757]

Model: "resnet"

Layer (type)	Output Shape	Param #	Connected to
img (InputLayer)	[(None, 784)]	0	
dense_35 (Dense)	(None, 128)	100480	img[0][0]
dense_36 (Dense)	(None, 128)	16512	dense_35[0][0]
dense_37 (Dense)	(None, 128)	16512	dense_36[0][0]
batch_normalization_20 (Batch Normalization)	(None, 128)	512	dense_37[0][0]
activation_20 (Activation)	(None, 128)	0	batch_normalization_20[0][0]
dropout_21 (Dropout)	(None, 128)	0	activation_20[0][0]
add_9 (Add)	(None, 128)	0	dropout_21[0][0] dense_36[0][0]
dense_38 (Dense)	(None, 128)	16512	add_9[0][0]
batch_normalization_21 (Batch Normalization)	(None, 128)	512	dense_38[0][0]
activation_21 (Activation)	(None, 128)	0	batch_normalization_21[0][0]
dropout_22 (Dropout)	(None, 128)	0	activation_21[0][0]
dense_39 (Dense)	(None, 128)	16512	dropout_22[0][0]
batch_normalization_22 (Batch Normalization)	(None, 128)	512	dense_39[0][0]
activation_22 (Activation)	(None, 128)	0	batch_normalization_22[0][0]
dropout_23 (Dropout)	(None, 128)	0	activation_22[0][0]
dense_40 (Dense)	(None, 128)	16512	dropout_23[0][0]
batch_normalization_23 (Batch Normalization)	(None, 128)	512	dense_40[0][0]
activation_23 (Activation)	(None, 128)	0	batch_normalization_23[0][0]

dropout_24 (Dropout)	(None, 128)	0	activation_23[0][0]
add_10 (Add)	(None, 128)	0	dropout_24[0][0] add_9[0][0]
dense_41 (Dense)	(None, 128)	16512	add_10[0][0]
dropout_25 (Dropout)	(None, 128)	0	dense_41[0][0]
dense_42 (Dense)	(None, 10)	1290	dropout_25[0][0]
=====			
Total params: 202,890			
Trainable params: 201,866			
Non-trainable params: 1,024			

Analysis:

- Training speed is improved compared to 10 resnets.
- Accuracy remained same. I think for MNIST data, we can get good accuracy with less number of trainable layers.

Step 8: Functional Model for CIFAR with the same baseline as MNIST.


```
In [21]: # Create the CIFAR baseline here.
# Create the functional Baseline here.
input_layer = keras.layers.Input(shape=cx_train.shape[1:])
il = Flatten()(input_layer)
h1 = keras.layers.Dense(64,activation="relu")(il)
h2 = keras.layers.Dense(64,activation="relu")(h1)
h3 = keras.layers.Dense(64,activation="relu")(h2)
output_layer = keras.layers.Dense(10, activation="softmax")(h3)
fcmodel = keras.models.Model(inputs=[input_layer], outputs=[output_layer])

fcmodel.compile(loss='categorical_crossentropy',
                optimizer=keras.optimizers.SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True),metrics=['accuracy'])

# Convert class vectors to binary class matrices.
cy_train = keras.utils.to_categorical(cy_train, 10)
cy_test = keras.utils.to_categorical(cy_test, 10)

fcmodel_detail = fcmodel.fit(cx_train, cy_train, epochs=20, batch_size=64, validation_split=0.1)

#plot accuracies for each epoch
fchistory = pd.DataFrame(fcmodel_detail.history)
fchistory.plot(figsize=(10,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.xlabel('Epochs')
plt.title('Functional Neural Nets')
plt.show()

#After training the model, evaluate the test set
fcmodel.evaluate(cx_test,cy_test)

#Print the summary of the model
fcmodel.summary()
```

Train on 45000 samples, validate on 5000 samples

Epoch 1/20

45000/45000 [=====] - 2s 49us/sample - loss: 2.0551 - accuracy: 0.2518 - val_loss: 2.0192 - val_accuracy: 0.2530

Epoch 2/20

45000/45000 [=====] - 2s 45us/sample - loss: 1.8717 - accuracy: 0.3299 - val_loss: 1.9179 - val_accuracy: 0.3186

Epoch 3/20

45000/45000 [=====] - 2s 45us/sample - loss: 1.7888 - accuracy: 0.3640 - val_loss: 1.7789 - val_accuracy: 0.3674

Epoch 4/20

45000/45000 [=====] - 2s 45us/sample - loss: 1.7383 - accuracy: 0.3819 - val_loss: 1.7761 - val_accuracy: 0.3690

Epoch 5/20

45000/45000 [=====] - 2s 44us/sample - loss: 1.7014 - accuracy: 0.3942 - val_loss: 1.7782 - val_accuracy: 0.3620

Epoch 6/20

45000/45000 [=====] - 2s 44us/sample - loss: 1.6642 - accuracy: 0.4099 - val_loss: 1.7122 - val_accuracy: 0.3922

Epoch 7/20

45000/45000 [=====] - 2s 45us/sample - loss: 1.6336 - accuracy: 0.4206 - val_loss: 1.6767 - val_accuracy: 0.4076

Epoch 8/20

45000/45000 [=====] - 2s 44us/sample - loss: 1.6056 - accuracy: 0.4280 - val_loss: 1.7612 - val_accuracy: 0.3822

Epoch 9/20

45000/45000 [=====] - 2s 44us/sample - loss: 1.5821 - accuracy: 0.4354 - val_loss: 1.7629 - val_accuracy: 0.3720

Epoch 10/20

45000/45000 [=====] - 2s 45us/sample - loss: 1.5592 - accuracy: 0.4449 - val_loss: 1.7059 - val_accuracy: 0.3902

Epoch 11/20

45000/45000 [=====] - 2s 45us/sample - loss: 1.5419 - accuracy: 0.4524 - val_loss: 1.6567 - val_accuracy: 0.4200

Epoch 12/20

45000/45000 [=====] - 2s 44us/sample - loss: 1.5217 - accuracy: 0.4564 - val_loss: 1.6910 - val_accuracy: 0.4006

Epoch 13/20

45000/45000 [=====] - 2s 44us/sample - loss: 1.5031 - accuracy: 0.4665 - val_loss: 1.6142 - val_accuracy: 0.4370

Epoch 14/20

45000/45000 [=====] - 2s 45us/sample - loss: 1.4863 - accuracy: 0.4712 - val_loss: 1.6621 - val_accuracy: 0.4128


```
Epoch 15/20
45000/45000 [=====] - 2s 44us/sample - loss: 1.4744 - accuracy: 0.4751 - val_loss:
1.6356 - val_accuracy: 0.4308
Epoch 16/20
45000/45000 [=====] - 2s 44us/sample - loss: 1.4618 - accuracy: 0.4801 - val_loss:
1.6317 - val_accuracy: 0.4328
Epoch 17/20
45000/45000 [=====] - 2s 45us/sample - loss: 1.4453 - accuracy: 0.4852 - val_loss:
1.6374 - val_accuracy: 0.4220
Epoch 18/20
45000/45000 [=====] - 2s 45us/sample - loss: 1.4319 - accuracy: 0.4888 - val_loss:
1.6045 - val_accuracy: 0.4340
Epoch 19/20
45000/45000 [=====] - 2s 45us/sample - loss: 1.4184 - accuracy: 0.4946 - val_loss:
1.6605 - val_accuracy: 0.4236
Epoch 20/20
45000/45000 [=====] - 2s 45us/sample - loss: 1.4029 - accuracy: 0.4992 - val_loss:
1.5867 - val_accuracy: 0.4498
```

Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x26dc0b11240>

Out[21]: (0, 1)

Out[21]: Text(0.5, 0, 'Epochs')

Out[21]: Text(0.5, 1.0, 'Functional Neural Nets')



10000/10000 [=====] - 0s 36us/sample - loss: 1.6179 - accuracy: 0.4300

Out[21]: [1.6179374866485596, 0.43]

Model: "model_6"

Layer (type)	Output Shape	Param #
=====		
input_7 (InputLayer)	[(None, 32, 32, 3)]	0
flatten_7 (Flatten)	(None, 3072)	0
dense_63 (Dense)	(None, 64)	196672
dense_64 (Dense)	(None, 64)	4160
dense_65 (Dense)	(None, 64)	4160
dense_66 (Dense)	(None, 10)	650
=====		
Total params: 205,642		
Trainable params: 205,642		
Non-trainable params: 0		
=====		

Analysis:

- Training speed is same for both the data datasets
- Accuracy is reduced significantly. I think, Cifar data has RGB layered images and the domain knowledge of the dataset is complex. Funtional neural network is not able to understand the complexity of the dataset which resulted in poor accuracies.

Step 9: Shallow ResNet for CIFAR with the same layers and compare it with MNIST.

```
In [55]: # Create the CIFAR ResNet here.
inputs = tf.keras.Input(shape=cx_train.shape[1:], name='img')
x = Dense(128, activation='relu')(inputs)
block_1_output = Dense(128, activation='relu')(x)

x = Dense(128)(block_1_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
y = Flatten()(x)
outputs = Dense(10, activation='softmax')(y)

rcmodel = tf.keras.Model(inputs, outputs, name='resnet')

rcmodel.compile(Adam(amsgrad=True), 'categorical_crossentropy', metrics=['accuracy'])

# Convert class vectors to binary class matrices.
cy_train = keras.utils.to_categorical(cy_train, 10)
cy_test = keras.utils.to_categorical(cy_test, 10)

rcmodel_detail = rcmodel.fit(cx_train, cy_train,
                             batch_size=128,
                             epochs=20,
                             validation_split=0.1)

#plot accuracies for each epoch
rchistory = pd.DataFrame(rcmodel_detail.history)
rchistory.plot(figsize=(10,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.xlabel('Epochs')
plt.title('Residual Neural Nets')
plt.show()

#After training the model, evaluate the test set
rcmodel.evaluate(cx_test,cy_test)
```

```
#Print the summary of the model  
rcmodel.summary()
```

Train on 45000 samples, validate on 5000 samples

Epoch 1/20

45000/45000 [=====] - 330s 7ms/sample - loss: 1.7810 - accuracy: 0.3831 - val_loss: 7.7729 - val_accuracy: 0.0976

Epoch 2/20

45000/45000 [=====] - 358s 8ms/sample - loss: 1.5499 - accuracy: 0.4672 - val_loss: 99.8459 - val_accuracy: 0.0976

Epoch 3/20

45000/45000 [=====] - 314s 7ms/sample - loss: 1.4754 - accuracy: 0.4922 - val_loss: 3827.4924 - val_accuracy: 0.0958

Epoch 4/20

45000/45000 [=====] - 314s 7ms/sample - loss: 17.9298 - accuracy: 0.3428 - val_loss: 102263.6198 - val_accuracy: 0.0950

Epoch 5/20

45000/45000 [=====] - 351s 8ms/sample - loss: 10.6031 - accuracy: 0.2106 - val_loss: 17954.0223 - val_accuracy: 0.0958

Epoch 6/20

45000/45000 [=====] - 363s 8ms/sample - loss: 6.3705 - accuracy: 0.2220 - val_loss: 11470.1549 - val_accuracy: 0.0958

Epoch 7/20

45000/45000 [=====] - 313s 7ms/sample - loss: 4.3110 - accuracy: 0.2392 - val_loss: 5624.3404 - val_accuracy: 0.0958

Epoch 8/20

45000/45000 [=====] - 339s 8ms/sample - loss: 3.4554 - accuracy: 0.2426 - val_loss: 3278.4047 - val_accuracy: 0.0958

Epoch 9/20

45000/45000 [=====] - 351s 8ms/sample - loss: 2.9452 - accuracy: 0.2556 - val_loss: 1714.8383 - val_accuracy: 0.0958

Epoch 10/20

45000/45000 [=====] - 349s 8ms/sample - loss: 2.4442 - accuracy: 0.2737 - val_loss: 767.6519 - val_accuracy: 0.0958

Epoch 11/20

45000/45000 [=====] - 313s 7ms/sample - loss: 2.2266 - accuracy: 0.2822 - val_loss: 1975.6464 - val_accuracy: 0.0950

Epoch 12/20

45000/45000 [=====] - 314s 7ms/sample - loss: 2.0512 - accuracy: 0.2875 - val_loss: 3693.7662 - val_accuracy: 0.0958

Epoch 13/20

45000/45000 [=====] - 311s 7ms/sample - loss: 1.9472 - accuracy: 0.3235 - val_loss: 88.7906 - val_accuracy: 0.0990

Epoch 14/20

45000/45000 [=====] - 564s 13ms/sample - loss: 1.9301 - accuracy: 0.3279 - val_loss: 5.8565 - val_accuracy: 0.0914



```
Epoch 15/20
45000/45000 [=====] - 336s 7ms/sample - loss: 1.9207 - accuracy: 0.3346 - val_loss:
16.4478 - val_accuracy: 0.0964
Epoch 16/20
45000/45000 [=====] - 323s 7ms/sample - loss: 1.9106 - accuracy: 0.3357 - val_loss:
246.8435 - val_accuracy: 0.1058
Epoch 17/20
45000/45000 [=====] - 326s 7ms/sample - loss: 1.9117 - accuracy: 0.3356 - val_loss:
291.4979 - val_accuracy: 0.1058
Epoch 18/20
45000/45000 [=====] - 324s 7ms/sample - loss: 1.9022 - accuracy: 0.3394 - val_loss:
3.0542 - val_accuracy: 0.1050
Epoch 19/20
45000/45000 [=====] - 318s 7ms/sample - loss: 1.9001 - accuracy: 0.3415 - val_loss:
35.9115 - val_accuracy: 0.1240
Epoch 20/20
45000/45000 [=====] - 325s 7ms/sample - loss: 1.8951 - accuracy: 0.3444 - val_loss:
24.2779 - val_accuracy: 0.0958
```

```
Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x26dfa32c0b8>
```

```
Out[55]: (0, 1)
```

```
Out[55]: Text(0.5, 0, 'Epochs')
```

```
Out[55]: Text(0.5, 1.0, 'Residual Neural Nets')
```



10000/10000 [=====] - 10s 1ms/sample - loss: 24.1797 - accuracy: 0.10003s - loss: 24.0904 - accuracy: 0 - ETA: 3s - loss: 24. - ETA: 2s - loss: 24.0850 - accuracy: 0.10003s - loss: 24.147

Out[55]: [24.179732403564454, 0.1]

Model: "resnet"

Layer (type)	Output Shape	Param #
img (InputLayer)	[(None, 32, 32, 3)]	0
dense_134 (Dense)	(None, 32, 32, 128)	512
dense_135 (Dense)	(None, 32, 32, 128)	16512
dense_136 (Dense)	(None, 32, 32, 128)	16512
batch_normalization_60 (Batch Normalization)	(None, 32, 32, 128)	512
activation_60 (Activation)	(None, 32, 32, 128)	0
dropout_63 (Dropout)	(None, 32, 32, 128)	0
dense_137 (Dense)	(None, 32, 32, 128)	16512
batch_normalization_61 (Batch Normalization)	(None, 32, 32, 128)	512
activation_61 (Activation)	(None, 32, 32, 128)	0
dropout_64 (Dropout)	(None, 32, 32, 128)	0
flatten_17 (Flatten)	(None, 131072)	0
dense_138 (Dense)	(None, 10)	1310730
Total params: 1,361,802		
Trainable params: 1,361,290		
Non-trainable params: 512		

Step 10: Custom ResNet.

Change the layer widths, dropout layers, batch sizes, and skip connections to see what we could do to make the ResNet better. Use the knowledge you gained from Lab 3 to do this.

```
In [50]: # Create the CIFAR ResNet here.
inputs = tf.keras.Input(shape=cx_train.shape[1:], name='img')
x = Dense(128, activation='relu')(inputs)
block_1_output = Dense(128, activation='relu')(x)

x = Dense(128)(block_1_output)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
x = Dense(128)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)
y = Flatten()(x)
outputs = Dense(10, activation='softmax')(y)
rc10model = tf.keras.Model(inputs, outputs, name='resnet')

rc10model.compile(optimizer=Adam(learning_rate=lr_schedule(0)), loss='categorical_crossentropy', metrics=['accuracy'])

# Convert class vectors to binary class matrices.
cy_train = keras.utils.to_categorical(cy_train, 10)
cy_test = keras.utils.to_categorical(cy_test, 10)

lr_scheduler = LearningRateScheduler(lr_schedule)

lr_reducer = ReduceLROnPlateau(factor=np.sqrt(0.1),
                                cooldown=0,
                                patience=5,
                                min_lr=0.5e-6)

callbacks = [lr_reducer, lr_scheduler]

rc10model_detail = rc10model.fit(cx_train, cy_train,
                                batch_size=32,
                                epochs=30,
                                validation_split=0.1,
                                shuffle=True,
                                callbacks=callbacks)

#plot accuracies for each epoch
rc10history = pd.DataFrame(rc10model_detail.history)
```

```
rc10history.plot(figsize=(10,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.xlabel('Epochs')
plt.title('Residual Neural Nets')
plt.show()

#After training the model, evaluate the test set
rc10model.evaluate(cx_test,cy_test)

#Print the summary of the model
rc10model.summary()
```

```
Learning rate: 0.001
Train on 45000 samples, validate on 5000 samples
Learning rate: 0.001
Epoch 1/30
45000/45000 [=====] - 298s 7ms/sample - loss: 1.8943 - accuracy: 0.3394 - val_loss:
596.7904 - val_accuracy: 0.0958
Learning rate: 0.001
Epoch 2/30
45000/45000 [=====] - 312s 7ms/sample - loss: 1.6885 - accuracy: 0.4205 - val_loss:
519.7225 - val_accuracy: 0.1070
Learning rate: 0.001
Epoch 3/30
45000/45000 [=====] - 302s 7ms/sample - loss: 1.6079 - accuracy: 0.4495 - val_loss:
343.1269 - val_accuracy: 0.0958
Learning rate: 0.001
Epoch 4/30
45000/45000 [=====] - 303s 7ms/sample - loss: 1.5752 - accuracy: 0.4606 - val_loss:
18261.2324 - val_accuracy: 0.0958
Learning rate: 0.001
Epoch 5/30
45000/45000 [=====] - 295s 7ms/sample - loss: 8.9512 - accuracy: 0.2420 - val_loss:
625.8508 - val_accuracy: 0.0958
Learning rate: 0.001
Epoch 6/30
45000/45000 [=====] - 296s 7ms/sample - loss: 1.9992 - accuracy: 0.2983 - val_loss:
1405.4320 - val_accuracy: 0.1038
Learning rate: 0.001
Epoch 7/30
45000/45000 [=====] - 296s 7ms/sample - loss: 1.9330 - accuracy: 0.3268 - val_loss:
928.6470 - val_accuracy: 0.0958
Learning rate: 0.001
Epoch 8/30
45000/45000 [=====] - 299s 7ms/sample - loss: 1.9090 - accuracy: 0.3355 - val_loss:
191.2962 - val_accuracy: 0.0958
Learning rate: 0.001
Epoch 9/30
45000/45000 [=====] - 297s 7ms/sample - loss: 1.9124 - accuracy: 0.3365 - val_loss:
3226.9604 - val_accuracy: 0.1038
Learning rate: 0.001
Epoch 10/30
45000/45000 [=====] - 305s 7ms/sample - loss: 1.8875 - accuracy: 0.3457 - val_loss:
1309.1224 - val_accuracy: 0.1050
Learning rate: 0.001
```

Epoch 11/30
45000/45000 [=====] - 326s 7ms/sample - loss: 1.8754 - accuracy: 0.3454 - val_loss:
1488.5722 - val_accuracy: 0.0940
Learning rate: 0.0001

Epoch 12/30
45000/45000 [=====] - 352s 8ms/sample - loss: 1.8525 - accuracy: 0.3561 - val_loss:
172.8564 - val_accuracy: 0.0984
Learning rate: 0.0001

Epoch 13/30
45000/45000 [=====] - 361s 8ms/sample - loss: 1.8458 - accuracy: 0.3571 - val_loss:
91.9831 - val_accuracy: 0.1462
Learning rate: 0.0001

Epoch 14/30
45000/45000 [=====] - 313s 7ms/sample - loss: 1.8418 - accuracy: 0.3612 - val_loss:
795.6470 - val_accuracy: 0.0986
Learning rate: 0.0001

Epoch 15/30
45000/45000 [=====] - 296s 7ms/sample - loss: 1.8421 - accuracy: 0.3583 - val_loss:
284.3303 - val_accuracy: 0.1164
Learning rate: 0.0001

Epoch 16/30
45000/45000 [=====] - 310s 7ms/sample - loss: 1.8282 - accuracy: 0.3646 - val_loss:
110.1915 - val_accuracy: 0.1152
Learning rate: 1e-05

Epoch 17/30
45000/45000 [=====] - 337s 7ms/sample - loss: 1.8205 - accuracy: 0.3701 - val_loss:
176.8868 - val_accuracy: 0.1066
Learning rate: 1e-05

Epoch 18/30
45000/45000 [=====] - 323s 7ms/sample - loss: 1.8184 - accuracy: 0.3675 - val_loss:
12.8034 - val_accuracy: 0.1542
Learning rate: 1e-05

Epoch 19/30
45000/45000 [=====] - 339s 8ms/sample - loss: 1.8143 - accuracy: 0.3686 - val_loss:
56.8887 - val_accuracy: 0.1138
Learning rate: 1e-05

Epoch 20/30
45000/45000 [=====] - 314s 7ms/sample - loss: 1.8162 - accuracy: 0.3681 - val_loss:
19.2211 - val_accuracy: 0.1254
Learning rate: 1e-05

Epoch 21/30
45000/45000 [=====] - 344s 8ms/sample - loss: 1.8177 - accuracy: 0.3675 - val_loss:
32.0102 - val_accuracy: 0.1648

```

Learning rate: 1e-06
Epoch 22/30
45000/45000 [=====] - 333s 7ms/sample - loss: 1.8171 - accuracy: 0.3693 - val_loss:
1.9610 - val_accuracy: 0.3392
Learning rate: 1e-06
Epoch 23/30
45000/45000 [=====] - 339s 8ms/sample - loss: 1.8119 - accuracy: 0.3729 - val_loss:
4.0643 - val_accuracy: 0.2242
Learning rate: 1e-06
Epoch 24/30
45000/45000 [=====] - 341s 8ms/sample - loss: 1.8122 - accuracy: 0.3703 - val_loss:
3.2406 - val_accuracy: 0.2410
Learning rate: 1e-06
Epoch 25/30
45000/45000 [=====] - 341s 8ms/sample - loss: 1.8136 - accuracy: 0.3705 - val_loss:
1.9577 - val_accuracy: 0.3368
Learning rate: 1e-06
Epoch 26/30
45000/45000 [=====] - 344s 8ms/sample - loss: 1.8200 - accuracy: 0.3684 - val_loss:
10.5099 - val_accuracy: 0.1650
Learning rate: 5e-07
Epoch 27/30
45000/45000 [=====] - 323s 7ms/sample - loss: 1.8193 - accuracy: 0.3720 - val_loss:
5.0268 - val_accuracy: 0.1936
Learning rate: 5e-07
Epoch 28/30
45000/45000 [=====] - 337s 7ms/sample - loss: 1.8104 - accuracy: 0.3718 - val_loss:
1.8643 - val_accuracy: 0.3548
Learning rate: 5e-07
Epoch 29/30
45000/45000 [=====] - 339s 8ms/sample - loss: 1.8162 - accuracy: 0.3705 - val_loss:
1.8718 - val_accuracy: 0.3552
Learning rate: 5e-07
Epoch 30/30
45000/45000 [=====] - 358s 8ms/sample - loss: 1.8195 - accuracy: 0.3702 - val_loss:
1.8797 - val_accuracy: 0.3546

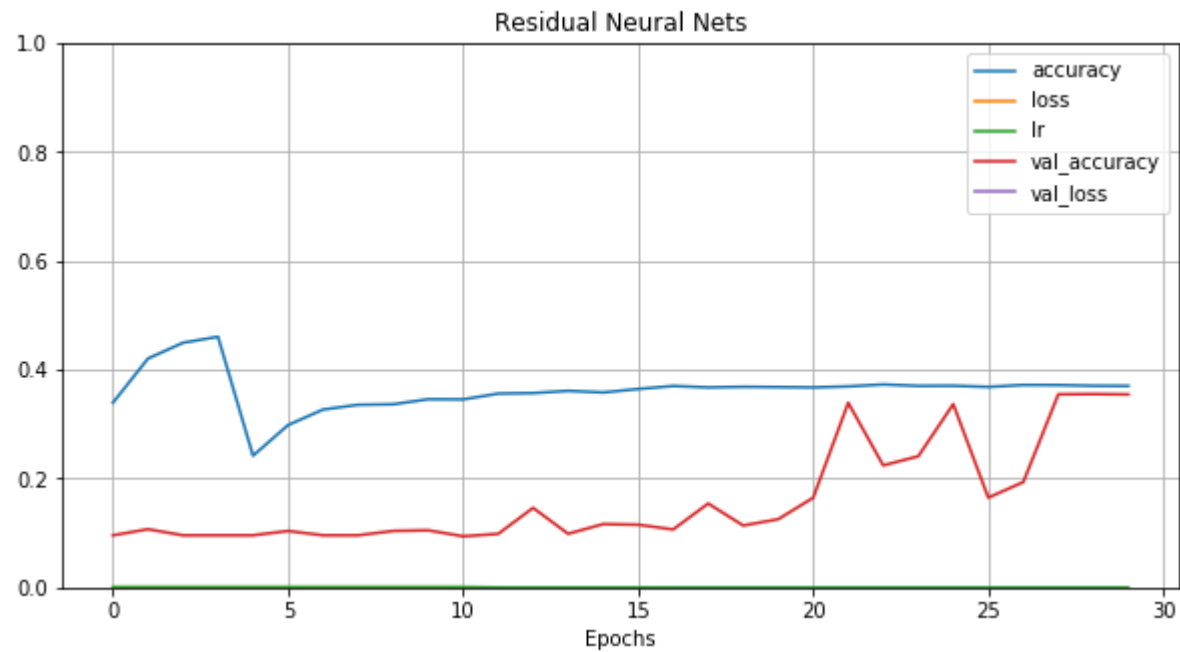
```

Out[50]: <matplotlib.axes._subplots.AxesSubplot at 0x26de6928e48>

Out[50]: (0, 1)

Out[50]: Text(0.5, 0, 'Epochs')

Out[50]: Text(0.5, 1.0, 'Residual Neural Nets')



10000/10000 [=====] - 13s 1ms/sample - loss: 1.8755 - accuracy: 0.3555

Out[50]: [1.875451368713379, 0.3555]

Model: "resnet"

Layer (type)	Output Shape	Param #
img (InputLayer)	[(None, 32, 32, 3)]	0
dense_124 (Dense)	(None, 32, 32, 128)	512
dense_125 (Dense)	(None, 32, 32, 128)	16512
dense_126 (Dense)	(None, 32, 32, 128)	16512
batch_normalization_56 (Batch Normalization)	(None, 32, 32, 128)	512
activation_56 (Activation)	(None, 32, 32, 128)	0
dropout_59 (Dropout)	(None, 32, 32, 128)	0
dense_127 (Dense)	(None, 32, 32, 128)	16512
batch_normalization_57 (Batch Normalization)	(None, 32, 32, 128)	512
activation_57 (Activation)	(None, 32, 32, 128)	0
dropout_60 (Dropout)	(None, 32, 32, 128)	0
flatten_15 (Flatten)	(None, 131072)	0
dense_128 (Dense)	(None, 10)	1310730
Total params: 1,361,802		
Trainable params: 1,361,290		
Non-trainable params: 512		

Analysis:

- We can see the performance of Residual nets are good compared to normal functional API. Accuracies are improved using Residual nets.
- By Increasing number trainable layers, training time is increasing significantly per epoch. But it might increase the accuracy slightly.
- I think, we need to use convolute neural net filters with proper hyper parameters to improve the accuracy.
- I tried, to tune the learning rate by using call back, but without proper filter, there is not increase in the accuracy.
- I think, we need to use convoluted neural nets with proper learning rate (Max learning rate/2). Max learning rate is the rate at which, accuracy declines.
- I tried batch sizes 128 and 32 but it did not impact the accuracy. I think, batch size would be subjective to other hyperparameters like learning rate and epochs.
- I tried, epochs 20 and 30. Both gave same performance. High value for epochs is not attempted because of high train speed. We can use early stopping call back to reduce training speed.

Step 11: Residual nets and gradient descent

- In a neural network, Gradient descent is used to minimize the error function and it is used during Forward propagation and Backward propagation.
- Forward propagation: Gradient descent is used to compute and pass the parameters to next layer. Parameters are calculated by minimize the cost function. Gradient descent is the method for minimization. In residual nets, residual layer also learns from the gradient and it will be added at the connection.
- Backward propagation: Gradient descent is used to update the parameters to train the network. In residual nets, residual layer are updated by the gradients and contribute in updating the parameters along the layer.
- Residual nets are used to pass the actual domain knowledge of input to next layers without manipulation, by skipping. In order to make the model better, we need to tune the hyperparameters of the module and more deeper layers should be added.

Step 12: Res nets in project

- Pros of using Res nets
 1. Our data has a certain level of complexity. (Not linearly separable). Training data is not huge. Using Res Nets, we can train better.
- Cons of using Res net
 1. As of data do not deal with images, number of input features is considerably less. As Residual nets have higher training speeds, compared to sequential and functional nets, we need to validate the performance of Functional API before using Residual nets. There is a chance that complexity of our data might not demand residual nets which would save the training speeds.

In []: